


Jaccard-constrained dense subgraph discovery

Chamalee Wickrama Arachchi, Nikolaj Tatti

University of Helsinki, Finland

DS23

What do we do?



- Three graph snapshots : G_1 , G_2 , and G_3
- Each graph snapshot contains 6 vertices, different edge sets.
- Three subgraphs : S_1 , S_2 , and S_3 .

Jaccard-constrained dense subgraph discovery problem

Input:

- A sequence of graph snapshots.
- Weight parameter : λ .

Output:

- A sequence of subgraphs.

Such that:

- $(\text{Sum of the densities} + \lambda \times \text{sum of pairwise similarities})$ is maximized.

Densest Subgraph Problem

- Definition of density:

$$d(S) = \frac{|E(S)|}{|V(S)|}$$

- Densest subgraph problem : Find a subset of vertices which maximizes the density.
- Exact¹ and greedy² algorithms.

¹ Andrew V Goldberg. Finding a maximum density subgraph. 1984.

² Moses Charikar. Greedy approximation algorithms for finding dense components in a graph. In International workshop on approximation algorithms for combinatorial optimization, pages 84–95. Springer, 2000.

Densest Common Subgraph Problem

- Extension for multiple graph snapshots.
- DCS : Find a **common** set of vertices which maximizes sum of the densities.
- **Variants of objectives :**
 - ▶ maximize sum of the densities ³.
 - ▶ maximize min density ⁴.

³Semertzidis, K., Pitoura, E., Terzi, E. and Tsaparas, P., 2019. Finding lasting dense subgraphs. *Data Mining and Knowledge Discovery*, 33(5), pp.1417-1445.

⁴Vinay Jethava and Niko Beerenwinkel. Finding dense subgraphs in relational graphs. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, pages 641–654. Springer, 2015.

Jaccard coefficient

- **Similarity** between two sets.
- S_i : subset of vertices from i th graph snapshot.
- S_j : subset of vertices from j th graph snapshot.

$$J(S_i, S_j) = \frac{|S_i \cap S_j|}{|S_i \cup S_j|}$$

- When $S_i = S_j$, pairwise Jaccard coefficient is one.

Definitions and our objective

- Density of i th subgraph

$$d(S_i) = \frac{|E(S_i)|}{|S_i|} .$$


- k : Length of the graph sequence
- Sum of the densities from all graph snapshots

$$d(\mathcal{S}) = \sum_{i=1}^k d(S_i) .$$

- Objective

$$q(\mathcal{S}; \lambda) = d(\mathcal{S}) + \lambda \sum_{i=1}^k \sum_{j=i+1}^k J(S_i, S_j) .$$

Example



- Let $\lambda = 0.3$, sets: $S_1 = \{a, b, d, f\}$, $S_2 = \{a, b, c, f\}$, and $S_3 = \{a, b, d, e, f\}$.
- The sum of densities is $\frac{4}{4} + \frac{6}{4} + \frac{8}{5} = 4.1$.
- $J(S_1, S_2) = \frac{3}{5} = 0.6$, $J(S_1, S_3) = \frac{4}{5} = 0.8$, and $J(S_2, S_3) = \frac{3}{6} = 0.5$.
- Our objective is equal to $4.1 + 0.3 \times (0.6 + 0.8 + 0.5) = 4.67$.

Our problem

- Given a graph sequence $= \{G_1, \dots, G_k\}$, with $G_i = (V, E_i)$, and a real number λ , find a collection of subset of vertices $\mathcal{S} = \{S_1, \dots, S_k\}$, such that $q(\mathcal{S})$ is maximized.
- It turns out that finding subgraphs maximizing our objective is **NP-hard**.

Algorithm - 1

Algorithm 1: ITR($\mathcal{G}, \lambda, \mathcal{S}$), finds subgraphs with good $q(\cdot; \lambda)$

```

1 while changes do
2   foreach  $i = 1, \dots, k$  do
3      $C \leftarrow V;$ 
4     foreach  $j = 2, \dots, |V|$  do
5        $u \leftarrow \arg \max_{v \in C} q(S_1, \dots, S_{i-1}, C \setminus \{v\}, S_{i+1}, \dots, S_k);$ 
6        $C \leftarrow C \setminus \{u\};$ 
7      $S_i \leftarrow$  best tested  $C$ , if the score improves;
8 return  $\mathcal{S};$ 

```

Figure: Iterative algorithm

- Initialization : DCS and individual densest subgraphs.

Algorithm - 2

Algorithm 2: GRD(\mathcal{G}, λ), finds subgraphs with good $q(\cdot; \lambda)$

- 1 $S \leftarrow S_1, \dots, S_k$, where $S_i = V$;
 - 2 **while** there are nodes **do**
 - 3 $u, j \leftarrow \arg \max_{v, i | v \in S_i} q(S_1, \dots, S_{i-1}, S_i \setminus \{v\}, S_{i+1}, \dots, S_k)$;
 - 4 $S_j \leftarrow S_j \setminus \{u\}$;
 - 5 **return** best tested S ;
-

Figure: Greedy algorithm

Experiments with Synthetic Datasets

- $|V^d|$ and $|V^s|$: Initial number of dense and sparse vertices.
- $|E|$: The expected number of edges.
- k : The number of snapshots.
- p_d , p_s , and p_c : The dense, sparse, and cross edge probabilities.

Dataset	$ V^d $	$ V^s $	$E[E]$	k	p_d	p_s	p_c	d_{true}	d_{dcs}	d_{ind}	J_{min}
Syn-1	100	900	672.8	10	0.05	0.0005	0.0005	34.68	24.26	35.74	0.47
Syn-2	100	5000	3850.2	5	0.05	0.0001	0.0001	40.47	13.03	40.57	0.18
Syn-3	120	1200	3922	5	0.06	0.005	0.002	27.4	17.62	27.62	0.45
Syn-4	250	5000	4709	8	0.03	0.001	0.001	55.79	29.65	55.98	0.27
Syn-5	500	3500	12136.29	7	0.05	0.0003	0.0003	112.11	87.27	112.12	0.53
Syn-6	350	3500	32015	5	0.06	0.005	0.002	67.96	52.17	67.96	0.49

- d_{dcs} : The density of densest common subgraph.
- d_{ind} : The sum of densities of locally densest subgraph from each graph snapshot.
- $J_{min} = \min_{i < j} J(V_i^d, V_j^d)$: The minimum Jaccard index between ground truth dense sets of vertices.
- d_{true} : The ground truth density of dense components.

- d_{dis} and q : The sum of densities and scores of the discovered sets.
- i : The number of iterations using Algorithm-1.
- ρ : The average Jaccard index between discovered and ground truth sets.
- $time$: The computational time in seconds.

Data	λ	ITR						GRD					
		d_{dis}	q	J_{min}	ρ	$time$	i	d_{dis}	q	J_{min}	ρ	$time$	
Syn-1	0.3	34.84	43.1	0.51	0.93	10	4	34.84	43.09	0.5	0.93	30	
	0.5	33.54	48.97	0.54	0.89	11	5	33.44	48.97	0.54	0.88	33	
	0.7	26.43	56.08	0.84	0.74	4	2	26.22	56.08	0.86	0.73	32	
Syn-2	0.3	40.57	41.18	0.18	0.99	50	3	40.57	40.77	0.18	0.99	92	
	0.4	40.57	41.39	0.18	0.99	53	3	40.57	41.39	0.18	0.99	84	
	0.5	40.57	41.6	0.18	0.99	50	3	40.57	41.6	0.18	0.99	87	
Syn-3	0.4	27.59	29.59	0.44	0.97	16	3	27.6	29.59	0.44	0.97	29	
	0.5	27.59	30.1	0.44	0.97	18	3	27.6	30.08	0.44	0.97	26	
	0.8	27.44	31.65	0.46	0.98	17	3	27.6	31.56	0.43	0.98	28	
Syn-4	0.4	55.75	60.38	0.27	0.98	66	2	55.89	60.42	0.27	0.98	307	
	0.5	55.82	61.56	0.27	0.98	132	4	55.83	61.55	0.27	0.98	308	
	0.6	55.76	62.71	0.27	0.98	132	4	55.79	62.71	0.27	0.98	285	
Syn-5	0.01	112.12	112.26	0.53	1	73	2	112.12	112.26	0.53	1	249	
	0.4	112.12	117.7	0.53	1	69	2	112.12	116.3	0.53	1	251	
	0.5	112.12	119.09	0.53	1	103	3	112.12	119.09	0.53	1	257	
Syn-6	0.1	67.96	68.61	0.49	1	42	2	67.96	68.61	0.49	1	226	
	0.8	67.96	73.18	0.49	1	67	3	67.96	73.17	0.49	1	252	

Experiments with Real-world Datasets

- $|V|$: The number of vertices.
- $|E|$: The number of edges.
- k : The number of snapshots.
- d_{dcs} : The density of densest common subgraph.
- d_{ind} : The sum of densities of locally densest subgraph from each graph snapshot.

Data	$ V $	$ E $	k	d_{dcs}	d_{ind}
Twitter-#	806	101.2	15	7.54	38.79
Enron	1 079	23.2	183	43.47	185.34
Facebook	4 117	83.13	104	14	88.65
Students	889	43.68	122	24.15	117.98
Twitter-user	4 605	109.19	93	23	90.63
Tumblr	1 980	65.3	89	36.67	103.98

Experiments with Real-world Datasets

- d_{dis} and q : The sum of densities and scores of the discovered sets.
- i : The number of iterations using Algorithm-1.
- $time$: The computational time in seconds.

Data	λ	ITR			i	GRD		
		d_{dis}	q	$time$		d_{dis}	q	$time$
Twitter-#	0.3	37.34	40.52	2.81	5	37.65	38.13	5.79
	0.5	14.14	54.54	3.06	7	31.01	42.37	6.38
	0.7	11	74.7	2.62	5	16.8	56.18	6.23
	0.8	11	83.8	2.79	6	9.33	75.47	6.88
Enron	0.05	130.43	357.29	72.02	5	122.08	343.02	816.54
	0.1	111.88	593.87	148.9	10	62.08	589.96	826.28
	0.5	95.88	2662.7	81.56	6	20	3139.25	854.43
	5	88.66	25254.81	98.87	7	20	31212.5	774.39
	0.1	52.86	126.98	214.74	7	56.03	85.64	9461.59
Facebook	0.5	36.4	446.09	148.94	6	25.47	344.6	10193.96
	0.7	31.25	657.92	185.91	7	23.41	475.83	9781.9
	1	27.41	922.45	136.72	6	22.48	659.32	9034.35
	0	108.9	108.9	124.38	4	102.7	102.7	2787.8
	0.2	46.99	526.19	100.17	5	33.17	567.84	2178.28
Students	0.5	46.17	1258.41	79.65	4	33.17	1369.86	2256.26
	0.8	44.5	2020.81	78.61	4	33.17	2171.88	2125.67
	0.01	79.21	81.61	572.25	8	63.53	68.53	7524.47
	0.1	12.04	269	260.31	9	13.9	212.28	6422.46
	0.2	11.49	545.93	149.91	5	11.64	423.1	6770.16
Twitter-user	0.5	11.49	1347.59	150.08	5	10.83	1040.05	6847.68
	0.1	71.33	316.07	50.23	4	66.47	302.83	1229.98
	0.5	64.37	1456.47	49.31	4	62.97	1291.82	1151.85
	0.7	59.25	2182.29	48.96	4	62.97	1783.35	1217.76

Conclusion

- Introduced a novel Jaccard weighted, dense subgraph discovery problem for graphs with multiple snapshots..
- Our goal was to find a dense subset of vertices from each graph snapshot such that the sum of densities and the similarity between the snapshots is maximized.
- Proved that our problem is **NP-hard**.
- Designed an iterative algorithm and a greedy algorithm which runs in polynomial time.
- Showed experimentally that the algorithms could find the ground truth using synthetic datasets.
- Showed experimentally that the number of iterations was low in iterative algorithm.
- Discovered dense collections in real world datasets and computational time is reasonable.

Thank you for your attention!!!

Speedup tricks

- The bottleneck in both algorithms : Finding the next vertex to delete.
- Select S_i , let $v \in S_i$ and $\mathcal{S}' := \mathcal{S}$ with S_i replaced with $S_i \setminus \{v\}$.
- The score difference between \mathcal{S} and \mathcal{S}' :

$$q(\mathcal{S}') - q(\mathcal{S}) = \frac{|E(S_i)| - \deg v}{|S_i| - 1} - \frac{|E(S_i)|}{|S_i|} + \lambda \sum_{j \neq i} J(S_i \setminus \{v\}, S_j) - J(S_i, S_j) \quad .$$

- In Alg-1 : To find the optimal v and i , we group the nodes in S_i such that the second term in objective is equal for the nodes in the same group.
- Maintain nodes in a collection of a priority queues keyed by its degree.
- To maintain the difference of the Jaccard indices, we maintain the sizes of intersection $|S_i \cap S_j|$ and the union $|S_i \cup S_j|$ for all i and j .
- To find the optimal v and i , we find the vertex with the smallest degree in each group, and then compare these candidates among different groups.